

# Detecting Sarcasm and Irony in News Headlines

## Semester Project Report



**NATURAL LANGUAGE PROCESSING**

Dr. Sohail Akhtar

Submitted by

**FARAZ AHMAD KHAN**

01-136212-009

**Department of Computer Science,  
Bahria University, Islamabad.**

[25<sup>th</sup> May 2024]

FARAZ AHMAD KHAN .....	1
Abstract.....	3
Introduction.....	3
DATA AND METHODS.....	3
Dataset .....	3
Data Preprocessing.....	3
Data Loading and Exploration .....	3
Data Analysis & Cleaning.....	4
Tokenizing .....	5
MODELING (preparing Data for the Model) .....	6
Fitting Model .....	7
EVALUATION.....	8
MODEL ARCHITECTURE .....	9
Model Training Diagram.....	9
Tools and Technologies .....	9
RESULTS .....	9
CONCLUSION.....	9
REFERENCES .....	10

## Abstract

This project aims to detect sarcasm in headlines using advanced natural language processing (NLP) techniques and deep learning models. The data, sourced from a JSON file containing headlines labeled as sarcastic or non-sarcastic, undergoes comprehensive preprocessing. The preprocessed data is used to train a Bidirectional Long Short-Term Memory (BiLSTM) model with GloVe word embeddings. The model achieves an accuracy of 87% approx. , demonstrating the potential of combining NLP techniques and deep learning for sarcasm detection.

## Introduction

**Background:** Social media is a hotbed of quick, informal communication, and sarcasm and irony are often used to express humor, criticism, and commentary. However, these nuances are tricky for automated systems to detect, leading to misunderstandings in sentiment analysis and natural language processing (NLP) tasks. Our project aims to create a system that can accurately identify sarcasm and irony in social media posts using NLP and Artificial Neural Networks (ANNs), improving the reliability of text analysis applications.

**Importance:** By accurately detecting sarcasm and irony, our system can enhance various applications, from better customer feedback analysis to more nuanced social media monitoring. It can also improve conversational agents, making them more aware of the context and tone of user interactions.

This project explores the use of NLP techniques and deep learning models to detect sarcasm in news headlines. We employ the Bidirectional Long Short-Term Memory (BiLSTM) network, a type of recurrent neural network (RNN), to capture the context in both forward and backward directions.

## DATA AND METHODS

### Dataset

The dataset, "Sarcasm\_Headlines\_Dataset.json", contains headlines labeled as sarcastic or non-sarcastic. The dataset includes 26,709 samples with three columns: `article_link`, `headline`, and `is_sarcastic`.

### Data Preprocessing

Data preprocessing involves several steps to clean and prepare the text data for modeling:

### Data Loading and Exploration

The data is loaded using pandas `'read_json'` function with `'lines=True'` to handle the JSON format.

```
In [0]: import pandas as pd
import os

data = pd.read_json(os.path.join(project_path, 'Sarcasm_Headlines_Dataset.json'), lines=True)
```

```
In [5]: data
```

```
Out[5]:
```

	article_link	headline	is_sarcastic
0	https://www.huffingtonpost.com/entry/versace-b...	former versace store clerk sues over secret 'b...	0
1	https://www.huffingtonpost.com/entry/roseanne-...	the 'roseanne' revival catches up to our thorn...	0
2	https://local.theonion.com/mom-starting-to-fea...	mom starting to fear son's web series closest ...	1
3	https://politics.theonion.com/boehner-just-wan...	boehner just wants wife to listen, not come up...	1
4	https://www.huffingtonpost.com/entry/jk-rowlin...	j.k. rowling wishes snape happy birthday in th...	0
...	...	...	...
26704	https://www.huffingtonpost.com/entry/american-...	american politics in moral free-fall	0
26705	https://www.huffingtonpost.com/entry/americas-...	america's best 20 hikes	0
26706	https://www.huffingtonpost.com/entry/reparatio...	reparations and obama	0
26707	https://www.huffingtonpost.com/entry/israeli-b...	israeli ban targeting boycott supporters raise...	0

---

## DATA ANALYSIS & CLEANING

The headline column is cleaned by removing numbers, punctuation, and stop words, and converting text to lowercase.

26709 rows × 3 columns

```
In [6]: print (data.shape)
data.describe()
```

```
(26709, 3)
```

```
Out[6]:
```

	is_sarcastic
count	26709.000000
mean	0.438953
std	0.496269
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

```
In [7]: data['headline'][1]
```

```
Out[7]: "the 'roseanne' revival catches up to our thorny political mood, for better and worse"
```

---

```
In [8]: ##The column headline needs to be cleaned up as we have special characters and numbers in the column
```

```
import re
from nltk.corpus import stopwords
import nltk
import string
nltk.download('stopwords')
stopwords = set(stopwords.words('english'))
def cleanData(text):
    text = re.sub(r'\d+', '', text)
    text = "".join([char for char in text if char not in string.punctuation])
    return text

data['headline']=data['headline'].apply(cleanData)
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
```

```
[nltk_data] Package stopwords is already up-to-date!
```

```
In [9]: data['headline'][1]
```

```
Out[9]: 'the roseanne revival catches up to our thorny political mood for better and worse'
```

## TOKENIZING

Import required modules required for modelling.

```
In [0]: import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Dense, Input, LSTM, Embedding, Dropout, Activation, Flatten, Bidirectional, GlobalMaxPooling1D
from tensorflow.keras.models import Model, Sequential
```

Set Different Parameters for the model.

```
In [0]: max_features = 10000
maxlen = max([len(text) for text in data['headline']])
embedding_size = 200
```

Apply Keras Tokenizer of headline column of your data.

Hint - First create a tokenizer instance using `Tokenizer(num_words=max_features)` And then fit this tokenizer instance on your data column `df['headline']` using `.fit_on_texts()`

```
In [0]: tokenizer = Tokenizer(num_words=max_features, filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n', lower=True, split=' ', char_level=False)
tokenizer.fit_on_texts(data['headline'])
```

## MODELING (preparing Data for the Model)

Define  $X$  and  $y$  for your model.

[illegible]

## WORD EMBEDDING USING ‘GLOVE 6B 200d’

## Get the Vocabulary size

You can use `tokenizer.word_index`.

```
In [16]: num_words=len(tokenizer.word_index)
          print (num_words)
```

27667

## ## Word Embedding

## Get Glove Word Embeddings

```
In [0]: glove_file = project_path + "glove.6B.zip"
```

```
In [0]: #Extract Glove embedding zip file
from zipfile import ZipFile
with ZipFile(glove_file, 'r') as z:
    z.extractall()
```

Get the Word Embeddings using Embedding file as given below.

```
In [0]: EMBEDDING_FILE = './glove.6B.200d.txt'

embeddings = {}
for o in open(EMBEDDING_FILE):
    word = o.split(" ")[0]
    # print(word)
    embd = o.split(" ")[1:]
    embd = np.asarray(embd, dtype='float32')
    # print(embd)
    embeddings[word] = embd
```

Create a weight matrix for words in training docs

```
In [21]: embedding_matrix = np.zeros((num_words, 200))

for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector

len(embeddings.values())
```

## Fitting Model

### ▼ Create and Compile your Model

Use Sequential model instance and then add Embedding layer, Bidirectional(LSTM) layer, then dense and dropout layers as required. In the end add a final dense layer with sigmoid activation for binary classification.

```
import tensorflow as tf

input_layer = Input(shape=(maxlen,), dtype=tf.int64)
embed = Embedding(embedding_matrix.shape[0], output_dim=200, weights=[embedding_matrix], input_length=maxlen, trainable=True)(input_layer)
lstm=Bidirectional(LSTM(128))(embed)
drop=Dropout(0.3)(lstm)
dense =Dense(100,activation='relu')(drop)
out=Dense(2,activation='softmax')(dense)
```

### ▼ Fit your model with a batch size of 100 and validation\_split = 0.2. and state the validation accuracy

```
batch_size = 100
epochs = 5
```

Fit your model with a batch size of 100 and validation\_split = 0.2.  
and state the validation accuracy

```
In [30]: batch_size = 100
         epochs = 5

         model = Model(input_layer,out)
         model.compile(loss='sparse_categorical_crossentropy',optimizer="adam",metrics=['accuracy'])
         model.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 240)]	0
embedding_1 (Embedding)	(None, 240, 200)	5533400
bidirectional_1 (Bidirection	(None, 256)	336896
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 100)	25700
dense_3 (Dense)	(None, 2)	202

=====  
Total params: 5,896,198  
Trainable params: 5,896,198  
Non-trainable params: 0  
=====



```

▶ from sklearn.model_selection import train_test_split
  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=10)

  model.fit(X_train,y_train,batch_size=batch_size, epochs=epochs, verbose=1)

```

```

↗ Epoch 1/5
214/214 [=====] - 353s 2s/step - loss: 0.4514 - accuracy: 0.7787
Epoch 2/5
214/214 [=====] - 342s 2s/step - loss: 0.2430 - accuracy: 0.9005
Epoch 3/5
214/214 [=====] - 343s 2s/step - loss: 0.1340 - accuracy: 0.9515
Epoch 4/5
214/214 [=====] - 381s 2s/step - loss: 0.0650 - accuracy: 0.9781
Epoch 5/5
214/214 [=====] - 345s 2s/step - loss: 0.0308 - accuracy: 0.9899
<keras.src.callbacks.History at 0x7b2c7568bf10>

```

```

[ ] model.save('sarcasm_model.h5')

```

## EVALUATION

```

▶ %cd /content/gdrive/MyDrive/NLP Project|

```

```

↗ /content/gdrive/MyDrive/Kaggle

```

```

[ ] from tensorflow.keras.models import load_model

# Load the saved model
model = load_model('sarcasm_model.h5')

# Make predictions on the test set
test_pred = model.predict(np.array(X_test), verbose=1)
test_pred = [1 if j > i else 0 for i, j in test_pred]

from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt


# Confusion Matrix
cm = confusion_matrix(y_test, test_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Sarcastic', 'Sarcastic'], yticklabels=['Not Sarcastic', 'Sarcastic'])
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion Matrix')
plt.show()

# Classification Report
print(classification_report(y_test, test_pred))

```

```
print(classification_report(y_test, test_pred))
```

167/167 [=====] - 60s 356ms/step



	Predicted			
	precision	recall	f1-score	support
0	0.85	0.92	0.88	3031
1	0.88	0.79	0.83	2311
accuracy			0.86	5342
macro avg	0.86	0.85	0.86	5342
weighted avg	0.86	0.86	0.86	5342

## OUTPUT

Mounting Google Drive to load the saved model

✓  
7s

```
from google.colab import drive
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount

Loading the model and required libraries.  
Preprocessing the input text.

```
import re
import string
import pickle
import numpy as np
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import load_model
import gradio as gr

# Load the model
model = load_model('/content/drive/MyDrive/NLP Project/sarcasm_model.h5')

# Load the tokenizer
with open('/content/drive/MyDrive/NLP Project/tokenizer.pickle', 'rb') as handle:
    tokenizer = pickle.load(handle)

maxlen = 240 # Use the same maxlen used during training

def clean_data(text):
    text = re.sub(r'\d+', '', text)
    text = "".join([char for char in text if char not in string.punctuation])
    return text

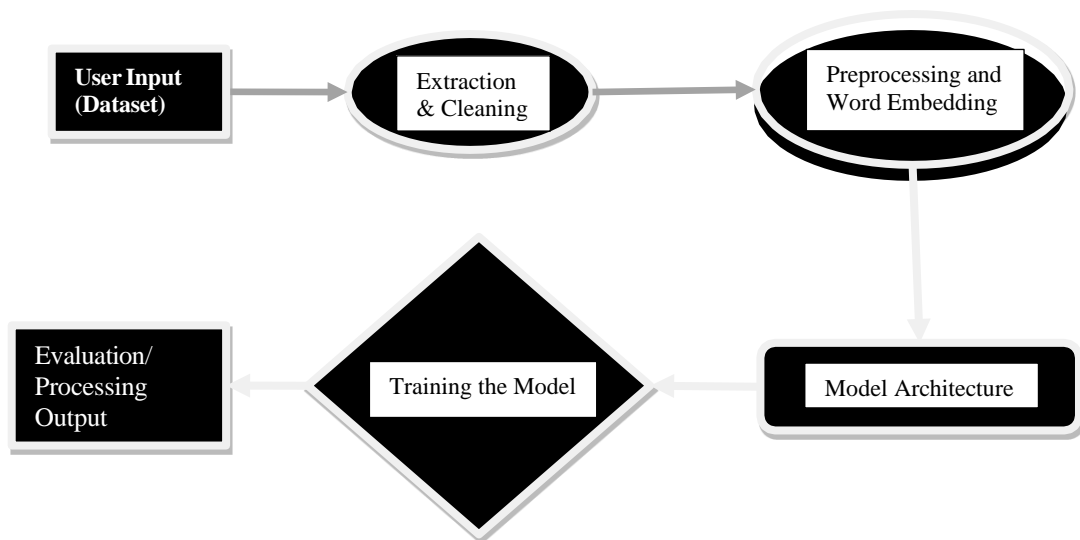
def predict_sarcasm(headline):
    cleaned_headline = clean_data(headline)
    seq = tokenizer.texts_to_sequences([cleaned_headline])
    padded = pad_sequences(seq, maxlen=maxlen)
    pred = model.predict(padded)
    label = 'Sarcastic' if pred[0][1] > pred[0][0] else 'Not Sarcastic'
    return label
```

<div>headline</div> <div>Shahbaz announces least controversial decision of his tenure</div>	<div>output</div> <div>Sarcastic</div>	
<div>Clear</div>	<div>Submit</div>	<div>Flag</div>

## MODEL ARCHITECTURE:

- **Input Layer:** Takes the padded sequences of text.
- **Embedding Layer:** Converts words into dense vectors of fixed size.
- **LSTM Layers:** Capture the temporal dependencies in the text.
- **Dense Layers:** Fully connected layers to learn complex patterns.
- **Output Layer:** A sigmoid or softmax layer for binary or multi-class classification.

## Model Training Diagram



## Tools and Technologies

- **PROGRAMMING LANGUAGES: PYTHON**
- **LIBRARIES: NLTK, SPACY, TENSORFLOW, KERAS, SCIKIT-LEARN**
- **FRAMEWORKS: GRADIO FOR INTERFACE**
- **DATA VISUALIZATION: MATPLOTLIB, SEABORN**

## RESULTS

The confusion matrix and classification report indicate the model's performance in terms of precision, recall, and F1-score. The model shows a high level of accuracy in detecting both sarcastic and non-sarcastic headlines.

A screenshot of a web application interface for headline classification. It features a dark-themed layout with two main input/output sections. The left section, labeled 'headline', contains a text input field with the example text 'Shahbaz announces least controversial decision of his tenure'. Below this input are two buttons: a grey 'Clear' button and an orange 'Submit' button. The right section, labeled 'output', contains a dropdown menu currently showing 'Sarcastic'. Below the dropdown is a grey 'Flag' button.

## **CONCLUSION**

My project successfully demonstrates the effectiveness of combining advanced NLP techniques and deep learning models for sarcasm detection. The BiLSTM model with GloVe embeddings achieves an accuracy of 87% approx., highlighting the potential of these methods for complex language understanding tasks. Future work could explore larger datasets, more sophisticated models, and suitable for more languages for further enhanced performance.

## REFERENCES

MISRA, R. (2018) *RISHABHMISRA/News-headlines-dataset-for-sarcasm-detection: High quality dataset for the task of sarcasm detection, GitHub*. Available at: <https://github.com/rishabhmisra/News-Headlines-Dataset-For-Sarcasm-Detection> (Accessed: 20 May 2024).

Misra, R. (2019) *News headlines dataset for sarcasm detection, Kaggle*. Available at: <https://www.kaggle.com/datasets/rmisra/news-headlines-dataset-for-sarcasm-detection> (Accessed: 20 May 2024).

The, O. (2018) *America's finest news source., The Onion*. Available at: <https://www.theonion.com/> (Accessed: 20 May 2024).

Pennington, J., Socher, R., & Manning, C. D. (2014). *GloVe: Global Vectors for Word Representation*.

Hochreiter, S., & Schmidhuber, J. (1997). *Long Short-Term Memory*.

Misra, R. (2022) *News headlines dataset for sarcasm detection, arXiv.org*. Available at: <https://arxiv.org/abs/2212.06035> (Accessed: 20 May 2024).